

Verlaxir: Verification of Message-Passing Systems

Matt Neave

Why must we verify message-passing systems?



Recent rise in cloud-based clusters



Distributed systems are hard to reason about




Message-passing systems are a common implementation of IPC



Lots of research into the verification of sequential / shared-memory solutions



Lacking research into verification tooling for message passing for actor-based systems




Concurrency

- Multiple computations executed at same time
- Threads / processes communicate with inter-process communication (IPC)
- IPC typically achieved with shared memory or message passing

Actor Model

- Model of computation where actors can:
 - Send messages to other actors
 - Spawn new actors
 - Act in response to a message

Actor Model



Elixir

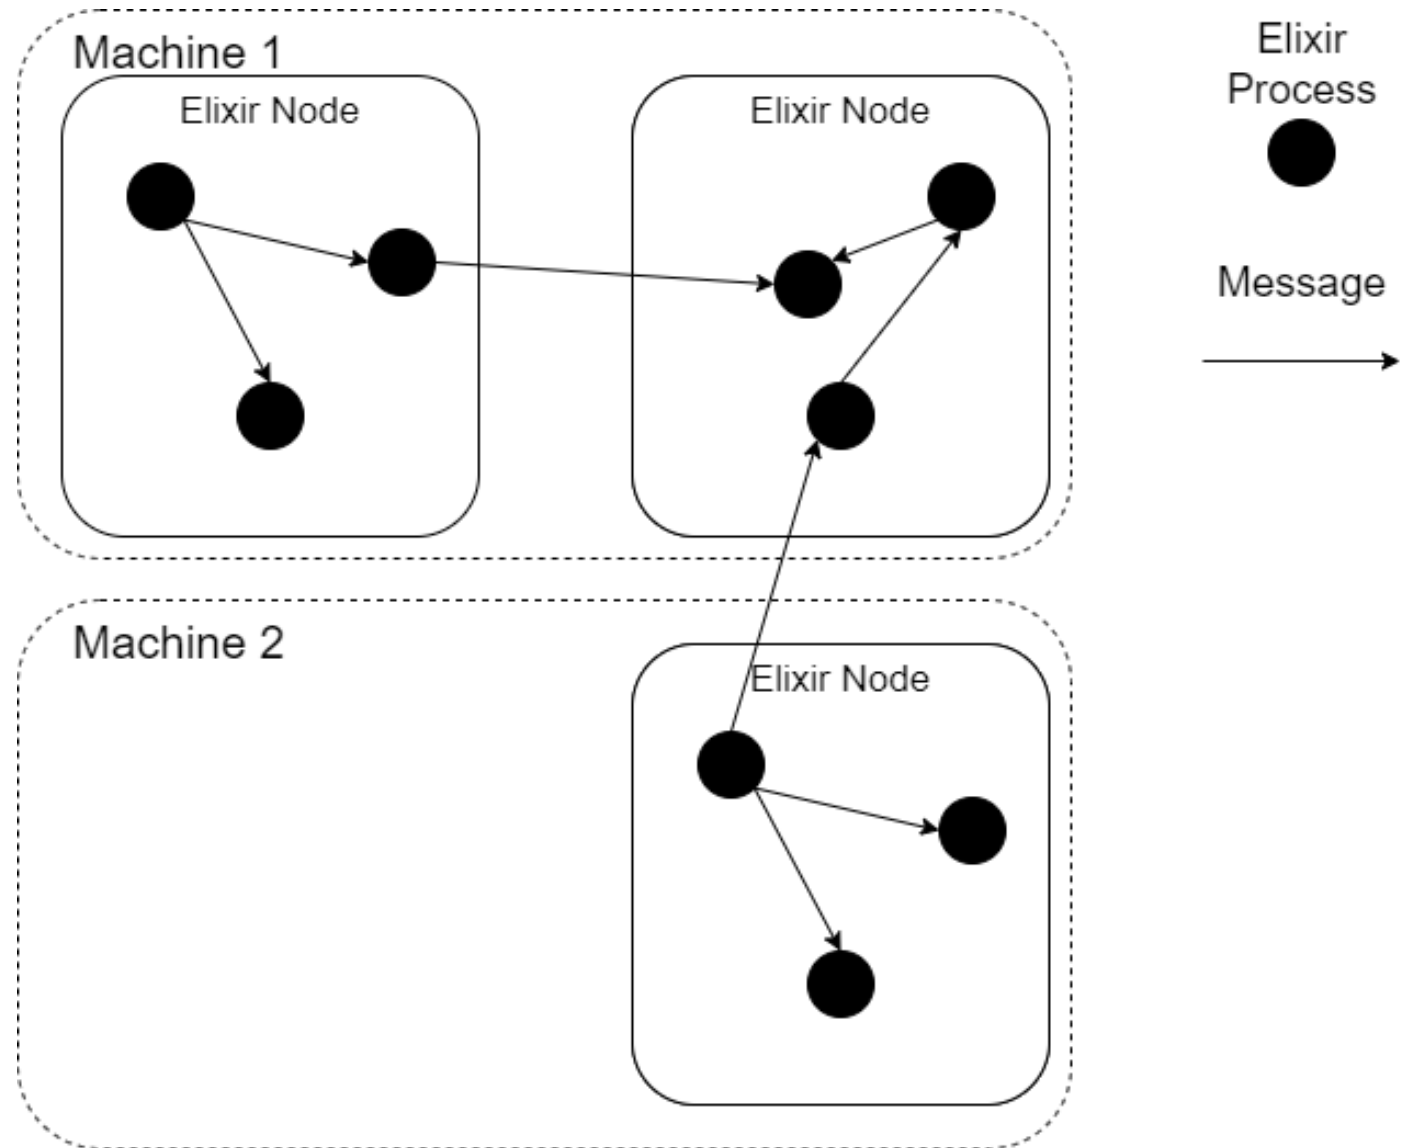
- Functional, concurrent, high-level programming language
- Builds on top of Erlang – sharing abstractions for distribution and fault-tolerance



Elixir

```
# Spawn a new process  
spawn(fn -> 1 + 2 end)
```

```
# Create a new BEAM instance  
Node.spawn(:"node1@localhost",  
  MyModule, :start, [])
```



Software Verification



Model Checking



Theorem Proving

Model Checking

Determining
whether a model
of system meets a
specification

Typically,
exhaustive search
of all states (prone
to state explosion)

Modern model
checkers include
Spin, PRISM,
BLAST...

Verification-aware Languages

- Programming language designed with built-in constructs to facilitate the verification of the program's correctness
- For example, Eiffel, **Dafny**

```
method sqrt(x: nat)
  requires x >= 0
  {
    ...
  }
```

Why must we verify message-passing systems?



Recent rise in cloud-based clusters



Distributed systems are hard to reason about



Message-passing systems are a common implementation of IPC



Lots of research into the verification of sequential / shared-memory solutions



Lacking research into verification tooling for message passing for actor-based systems

Verlixir

- Verification-aware programming language for message-passing systems
- Supports the simulation and verification of systems in a safe, controlled environment
- Provides a specification language with support for linear temporal formulae, predicates and design-by-contract-style functions
- Allows concurrent properties to be parameterized and verified over multiple configurations

Verlixir Simulator



Verlixir Simulator

```
defmodule Coordinator do
```

```
  @init
```

```
  @spec start_coordinator() :: :ok
```

```
  def start_coordinator do
```

```
    client_count = 100
```

```
    server_count = 10
```

```
    ... spawn and link servers / clients ...
```

```
  end
```

```
end
```

Entry point for Verlixir



Verlixir Simulator

`./verlixir -s distributed_braodcast.ex`

Starting global broadcast...

South-America broadcasting...

EU-South broadcasting...

...

All client acknowledgements received

Because of delays, real-world system may not observe this interleaving for hours / days / weeks...






Verlixir Verifier - Making Elixir Verification-aware

- Deadlock and livelock detection, for FREE
- Linear temporal logic (LTL) to reason about time and change
- Design-by-Contract functions using pre- and post-conditions
- Predicates


Verlaxir Verifier: Deadlocks

- Dining philosophers – classical deadlock problem
- Philosophers need two forks to eat!



Verlixir Verifier: Deadlocks

- Dining philosophers – classical deadlock problem
- Philosophers need two forks to eat!



Some simulations output:

Picking up 1
Picking up 2

Picking up 2
Picking up 1

All philosophers have finished eating!

Others output:

Picking up 1
Picking up 2

Picking up 1
Picking up 2

timeout

Verlir Verifier: Counterexamples

- Verlir verification mode produces counterexample!

```
verlir -v dining_philosophers.ex
```

Model ran successfully. 1 error(s) found.


The program likely reached a deadlock. Generating trace.

```
(proc_1 - fork_loop) recv [PICKUP, phil1]  
(proc_1 - fork_loop) send [ACK, phil1]  
(proc_2 - fork_loop) recv [PICKUP, phil2]  
(proc_2 - fork_loop) send [ACK, phil2]
```

```
(proc_2 - fork_loop) recv [PICKUP, phil1]  
(proc_1 - fork_loop) recv [PICKUP, phil2]
```

```
timeout
```

Both forks have been
picked up!



Verlir Verifier: LTL

- Same as deadlock counterexamples, Verlir will produce counterexamples violating LTL formula

```
defmodule Atm do

  @ltl """
  <>(card_inserted)
  [] (card_inserted -> (<> [] (transaction_failed) || <> [] (atm_balance < 1000)))
  """

  def start_atm do
    atm_balance = 1000
    card_inserted = false
    transaction_failed = false
    ... implementation ...
  end
end
```

Verlixir Verifier: contracts

- Define pre- and post-condition on contract

```
defv add(x, y), pre: x > 0 && y > 0, post: z == x + y do
  ... defines z somehow ...
end
```

↖
`defv` introduces a contract

Verlir: predicates

- Simplify LTL claim readability
- Used as expression conditions

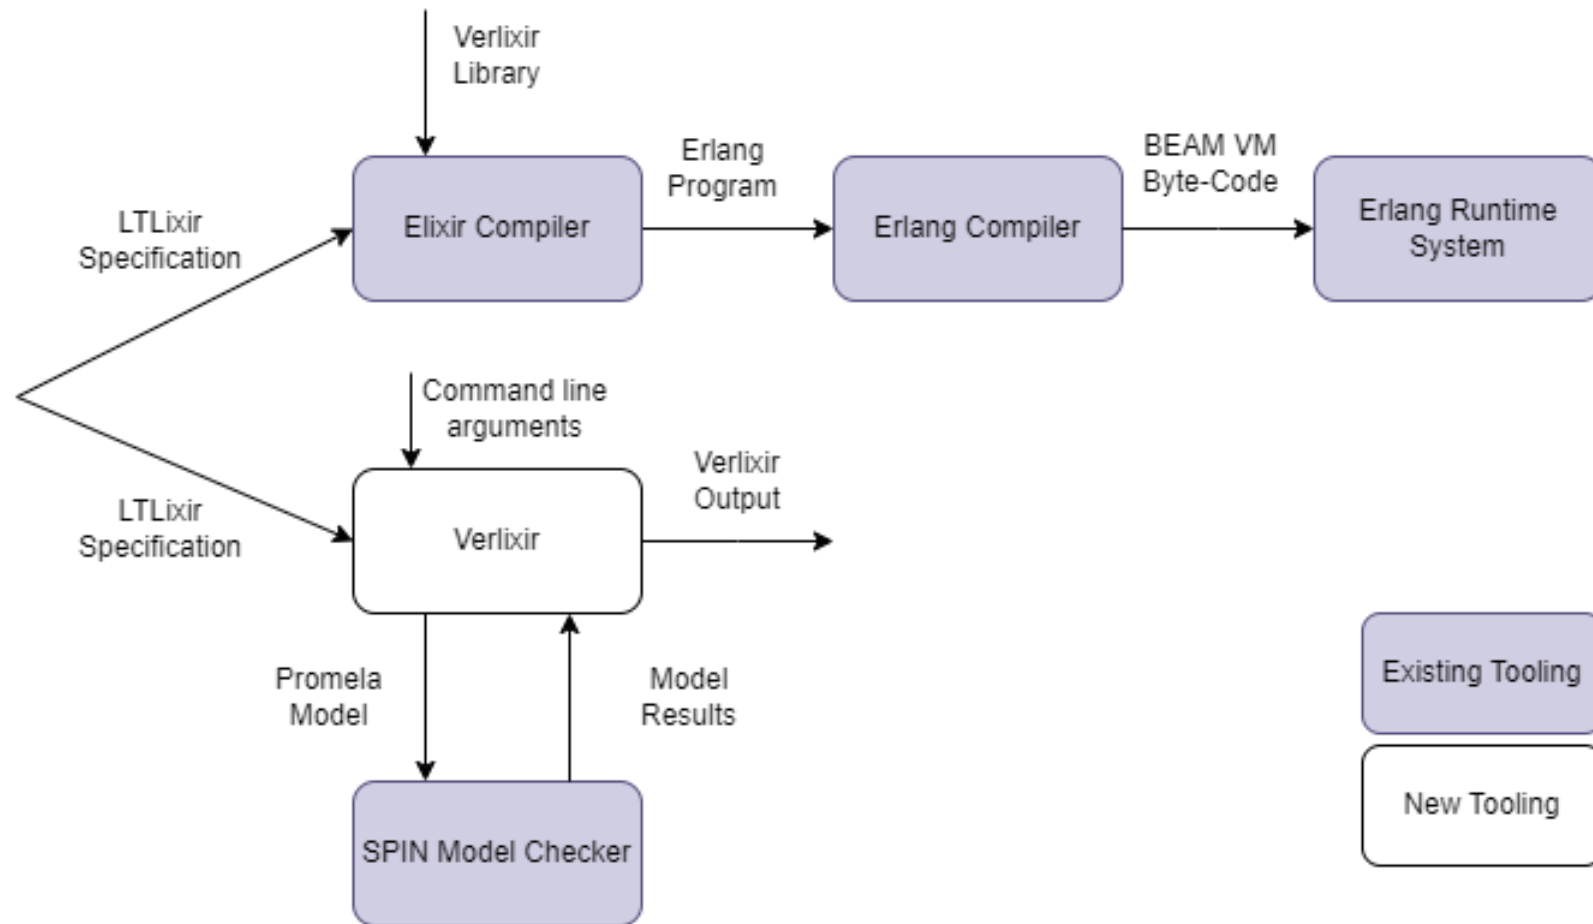
```
@ltl "(new_round -> <>(consensus))"
def run do
  predicate new_round, messages_received == 0 && messages_sent == 0
  predicate consensus, messages_received >= number_of_nodes / 2 + 1
  ...
  if new_round do
    ...
  end
end
```

Verlixir Verifier: concurrency parameterisation

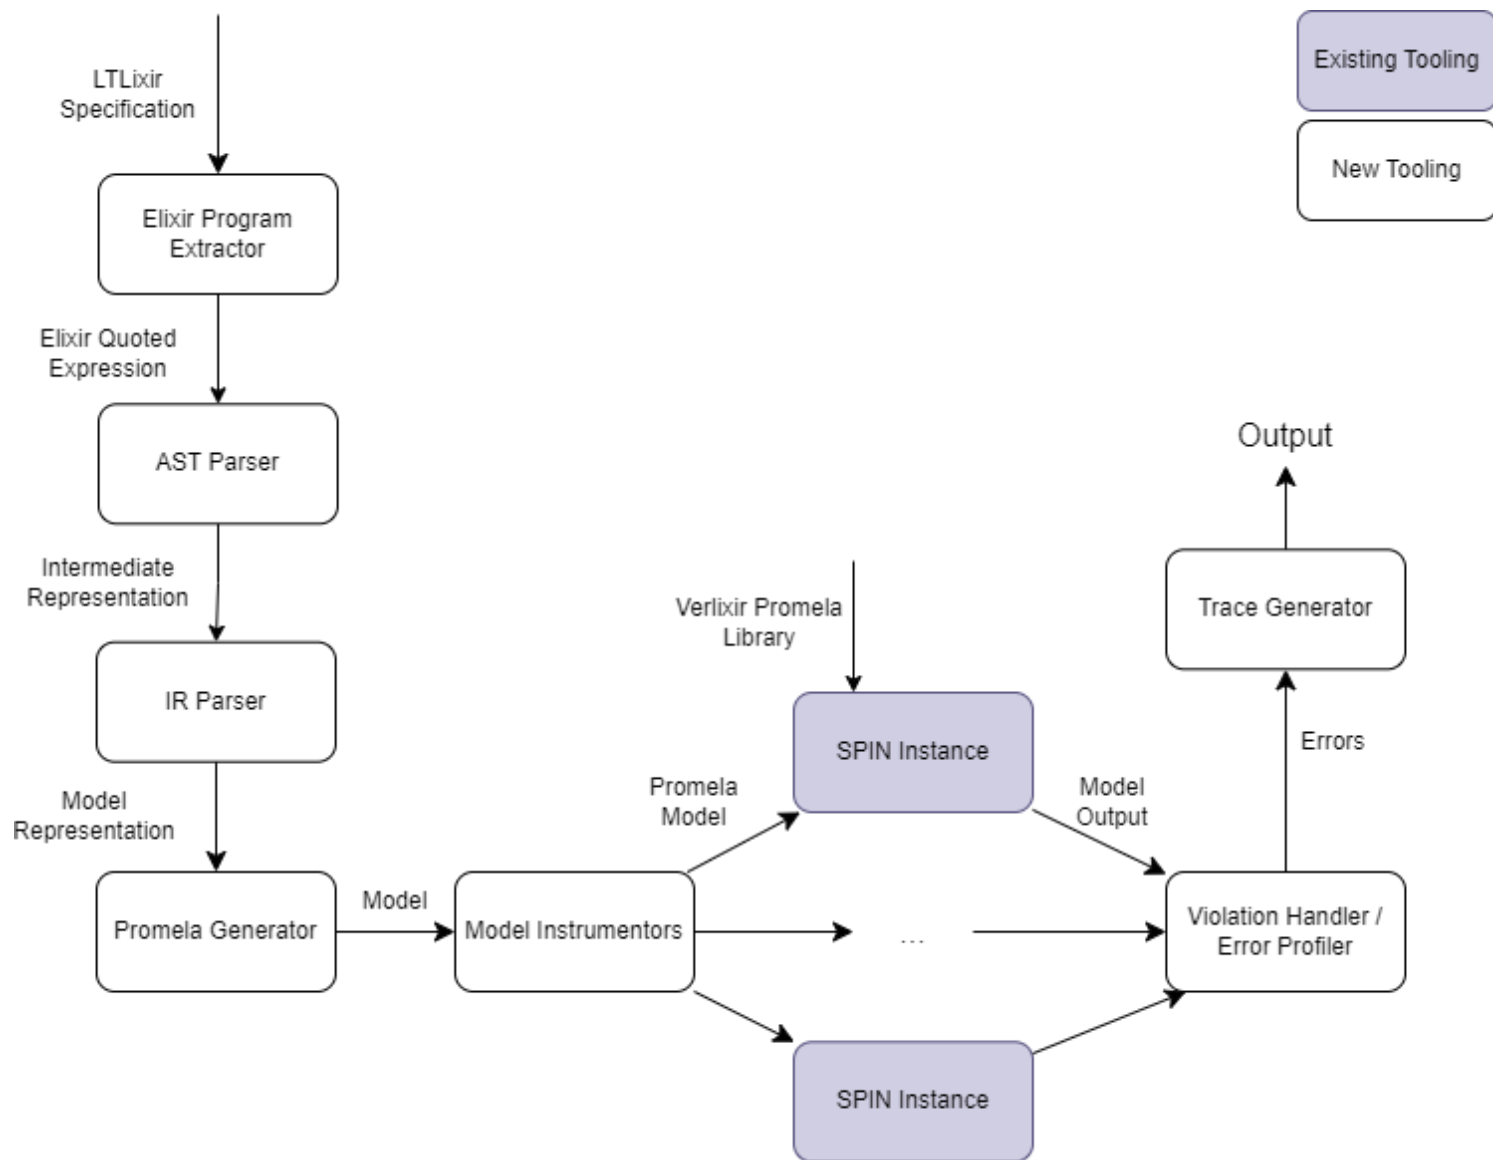
- Model different configurations of variables

```
@model { :client_count, :server_count }  
def start_coordinator do  
  client_count = 100  
  server_count = 10  
  ... spawn and link servers / clients ...  
end
```

High-level Overview



System Overview



Evaluated algorithms

- Basic paxos
- Consistent hash ring
- Dining philosophers
- Two-phase commit
- Raft (leader election)

Future Work

- Combine model checking with theorem proving to build a more robust verification tool
- Fault injection – determine if systems are fault tolerant by killing processes during verification
- Computation-tree logic (CTL) support, i.e. for algorithms like Raft, where leader election is not guaranteed, we can determine there exists a path where a leader is elected

Final Remarks

- Specifying systems in terms of safety and liveness should be apart of all software development
- Verification-aware languages make system specification more accessible to developers
- With rise of generative AI, we could move towards automatic system implementation from well formed system properties – the use of verification-aware languages can form a feedback loop to ensure the system correctness



Questions?